

# SQL SATURDAY - SALT LAKE CITY - 2020

## SQL LOGIN DATA WAREHOUSE

GEORGE WALKEY

SR SQL DBA

PATIENTFIRST CORP

RICHMOND, VA

[HTTPS://GITHUB.COM/GWALKEY](https://github.com/gwalkey)

# ANSWER MY LOGIN QUESTION:



Who



How Often



Server



Database



Client/Host



Application



When





# PIECES OF THE SOLUTION

- Extended Event Sessions
- Powershell XE Importer
- OLTP Relational Database
- OLAP Star Schema Database
- SSAS Tabular Database
- SQL Agent Jobs – Data Collectors/Aggregators
- Excel Pivot Table as Front-End

# EXTENDED EVENT SESSION

- Collects only the XE Login Event to file
- Collects only successful logins
- XE is very lightweight
- Even with fast-polling apps (50 logins/sec)
- XE Sessions automatically start when SQL Starts
- Logging 9M+ Events a day for one server
- The XEL trace files are copied up to a Central Server for ETL and reporting

# POWERSHELL EXTENDED EVENTS IMPORTER

- MS Recommends
  - <https://docs.microsoft.com/en-us/sql/relational-databases/system-functions/sys-fn-xe-file-target-read-file-transact-sql?view=sql-server-2017>
  - 11 Hours on 9M+ Trace file
- Dan Guzman – SQL MVP – (**QueryableXEventData**)
  - <https://dba.stackexchange.com/questions/206863/what-is-the-right-tool-to-process-big-xel-files-sql-server-extended-events-log>
  - **11 Minutes on 9M+ Trace file**
- Load 2 Standard SQL Server Assemblies
  - C:\Program Files\Microsoft SQL Server\150\Shared\Microsoft.SqlServer.XEvent.Linq.dll
  - C:\Program Files\Microsoft SQL Server\150\Shared\Microsoft.SqlServer.XE.Core.dll

# OLTP RELATIONAL DATABASE

- Used to Bulk-Load the XE session files from all Servers into SQL
  - Aggregate these Metrics:
    - Server
    - Database
    - Host
    - Application String
    - User/Login
    - Hour of Day
    - **Connection Count per Hour**
  - ETL Errors go to Errors Table for inspection during QA Pass



# OLAP STAR SCHEMA DATABASE

- Typical Star-Schema Database
  - Used to feed the SSAS Tabular Model
    - Date Dimension
    - Server Dimension
    - Database Dimension
    - Host Dimension
    - Application Dimension
    - Login/User Dimension
    - Single Fact Table Measure: Logins per Hour

# ORIGINAL WAS SSAS MD CUBE

- Why Multidimensional?
  - Project Started in 2014
  - My First Data Warehouse
  - Tabular still 1.0-ish in 2014
  - Didn't have enough \$\$\$ Memory on Server for Tabular
  - Didn't have Enterprise Edition
  - Learned Data Warehousing using MD and MDX
- Not showing you how to build the Cube today
  - Takes too Long
  - Download the ABF and Restore it from my Github repo
  - MOLAP Cube (ROLAP/Clustered ColumnStore an Option)



# SQL AGENT JOBS

- Jobs on each Monitored Server
  - XE Session Target Filename changes at Midnight
  - XE Session Target Filename matches Day of Week
- Jobs On Central Database
  - Copies up XE Session files from above servers
  - Import and Clean all XE Session Files
  - ETL loads OLTP and OLAP Databases
  - Process Tabular Model

# EXCEL FRONT-END

- Why Excel and Not Power BI?
  - Pivot Tables fit better
  - Easier Drill-Down
  - Easier Filtering

# OPERATIONAL STATISTICS

- Running since 2015
  - 20 Servers
  - 275 Databases
  - 7500 Hosts
  - 680 Application Strings
  - 8300 User/Logins
  - Fact Table 30M+ Rows
  - 15 Billion Logins
- Cube is only 350MB
- Entire Daily Import Process takes 16 Minutes



The background is a deep blue to purple gradient. It features several abstract elements: a large circular scale on the left with markings from 140 to 260; several concentric circles and arcs in white and light blue; and a dense field of out-of-focus circular bokeh in shades of blue and purple.

DEMO

[HTTPS://GITHUB.COM/GWALKEY](https://github.com/gwalkey)



THANK YOU!

[HTTPS://GITHUB.COM/GWALKEY](https://github.com/gwalkey)